

Keras -- MLPs on MNIST

```
In [1]: # if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow"
        use this command

from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal

#install the nightly artifacts of tensorflow in order to use TensorFlow eager mode.
#It is a new, experimental feature that is not yet included in the releases.
```

Using TensorFlow backend.

```
C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:526: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
```

```
C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:527: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
```

```
C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:528: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
```

```
C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:529: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
```

```
C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:530: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
_np_qint32 = np.dtype [("qint32", np.int32, 1)]
```

```
C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\framework\dtypes.py:535: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,)type'.
```

```
np_resource = np.dtype [("resource", np.ubyte, 1)]
```

```
In [2]: %matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

```
In [3]: # the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

```
In [4]: print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

```
In [5]: # if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [6]: # after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
In [7]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  45 186 253 253 150  27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  16  93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  249 253 249  64  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  24 114 221 253 253 253
253 201  78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  23  66 213 253 253 253 253 198  81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  18 171 219 253 253 253 253 195
 80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 55 172 226 253 253 253 253 244 133  11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132  16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

```
In [8]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

X_train = X_train/255
X_test = X_test/255
```

```
In [9]: # example data point after normlizing  
print(X_train[0])
```

[0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.01176471	0.07058824	0.07058824	0.07058824
0.49411765	0.53333333	0.68627451	0.10196078	0.65098039	1.
0.96862745	0.49803922	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.11764706	0.14117647	0.36862745	0.60392157
0.66666667	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.88235294	0.6745098	0.99215686	0.94901961	0.76470588	0.25098039
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.19215686
0.93333333	0.99215686	0.99215686	0.99215686	0.99215686	0.99215686
0.99215686	0.99215686	0.99215686	0.98431373	0.36470588	0.32156863
0.32156863	0.21960784	0.15294118	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.07058824	0.85882353	0.99215686
0.99215686	0.99215686	0.99215686	0.99215686	0.77647059	0.71372549
0.96862745	0.94509804	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.31372549	0.61176471	0.41960784	0.99215686
0.99215686	0.80392157	0.04313725	0.	0.16862745	0.60392157
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.05490196	0.00392157	0.60392157	0.99215686	0.35294118
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.54509804	0.99215686	0.74509804	0.00784314	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.04313725
0.74509804	0.99215686	0.2745098	0.	0.	0.

0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.1372549	0.94509804
0.88235294	0.62745098	0.42352941	0.00392157	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.31764706	0.94117647	0.99215686
0.99215686	0.46666667	0.09803922	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.17647059	0.72941176	0.99215686	0.99215686
0.58823529	0.10588235	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.0627451	0.36470588	0.98823529	0.99215686	0.73333333
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.97647059	0.99215686	0.97647059	0.25098039	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.18039216	0.50980392	0.71764706	0.99215686
0.99215686	0.81176471	0.00784314	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.15294118	0.58039216
0.89803922	0.99215686	0.99215686	0.99215686	0.98039216	0.71372549
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.09411765	0.44705882	0.86666667	0.99215686	0.99215686	0.99215686
0.99215686	0.78823529	0.30588235	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.09019608	0.25882353	0.83529412	0.99215686
0.99215686	0.99215686	0.99215686	0.77647059	0.31764706	0.00784314
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.07058824	0.67058824
0.85882353	0.99215686	0.99215686	0.99215686	0.99215686	0.76470588
0.31372549	0.03529412	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.21568627	0.6745098	0.88627451	0.99215686	0.99215686	0.99215686
0.99215686	0.95686275	0.52156863	0.04313725	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.53333333	0.99215686
0.99215686	0.99215686	0.83137255	0.52941176	0.51764706	0.0627451
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.
0.	0.	0.	0.	0.	0.

[illegible]

```
In [10]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
```

Softmax classifier

```

In [11]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias)
# where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```



```
In [12]: # some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

```
In [13]: # start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))
```

WARNING:tensorflow:From C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\ops\resource_variable_ops.py:435: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:
Colocations handled automatically by placer.

```

In [14]: # Before training a model, you need to configure the learning process, which is done
          # one via the compile method

          # It receives three arguments:
          # An optimizer. This could be the string identifier of an existing optimizer, https://keras.io/optimizers/
          # A loss function. This is the objective that the model will try to minimize., https://keras.io/losses/
          # A list of metrics. For any classification problem you will want to set this to
            metrics=['accuracy']. https://keras.io/metrics/

          # Note: when using the categorical_crossentropy loss, your targets should be in categorical format
          # (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vector that is all-zeros except
          # for a 1 at the index corresponding to the class of the sample).

          # that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None,
# validation_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoch=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss values and
# metrics values at successive epochs, as well as validation loss values and validation metrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```

WARNING:tensorflow:From C:\Users\wwang26\AppData\Local\Continuum\anaconda3\lib\site-packages\tensorflow\python\ops\math_ops.py:3066: to_int32 (from tensorflow.python.ops.math_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.cast instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 2s 35us/step - loss: 1.2756 - accuracy: 0.6960 - val_loss: 0.8088 - val_accuracy: 0.8340

Epoch 2/20

60000/60000 [=====] - 2s 34us/step - loss: 0.7148 - accuracy: 0.8408 - val_loss: 0.6054 - val_accuracy: 0.8616

Epoch 3/20

60000/60000 [=====] - 3s 44us/step - loss: 0.5861 - accuracy: 0.8606 - val_loss: 0.5241 - val_accuracy: 0.8731

Epoch 4/20

60000/60000 [=====] - 2s 41us/step - loss: 0.5246 - accuracy: 0.8695 - val_loss: 0.4787 - val_accuracy: 0.8796

Epoch 5/20

60000/60000 [=====] - 2s 35us/step - loss: 0.4871 - accuracy: 0.8763 - val_loss: 0.4485 - val_accuracy: 0.8862

Epoch 6/20

60000/60000 [=====] - 2s 41us/step - loss: 0.4613 - accuracy: 0.8805 - val_loss: 0.4277 - val_accuracy: 0.8888

Epoch 7/20

60000/60000 [=====] - 2s 30us/step - loss: 0.4422 - accuracy: 0.8841 - val_loss: 0.4117 - val_accuracy: 0.8920

Epoch 8/20

60000/60000 [=====] - 2s 26us/step - loss: 0.4273 - accuracy: 0.8870 - val_loss: 0.3988 - val_accuracy: 0.8950

Epoch 9/20

60000/60000 [=====] - 2s 26us/step - loss: 0.4153 - accuracy: 0.8893 - val_loss: 0.3883 - val_accuracy: 0.8972

Epoch 10/20

60000/60000 [=====] - 2s 25us/step - loss: 0.4053 - accuracy: 0.8913 - val_loss: 0.3795 - val_accuracy: 0.8988

Epoch 11/20

60000/60000 [=====] - 1s 24us/step - loss: 0.3969 - accuracy: 0.8929 - val_loss: 0.3723 - val_accuracy: 0.9005

Epoch 12/20

60000/60000 [=====] - 2s 36us/step - loss: 0.3896 - accuracy: 0.8948 - val_loss: 0.3659 - val_accuracy: 0.9013

Epoch 13/20

60000/60000 [=====] - 2s 34us/step - loss: 0.3832 - accuracy: 0.8959 - val_loss: 0.3603 - val_accuracy: 0.9033

Epoch 14/20

60000/60000 [=====] - 2s 35us/step - loss: 0.3776 - accuracy: 0.8974 - val_loss: 0.3554 - val_accuracy: 0.9047

Epoch 15/20

60000/60000 [=====] - 2s 32us/step - loss: 0.3726 - accuracy: 0.8982 - val_loss: 0.3511 - val_accuracy: 0.9055

Epoch 16/20

60000/60000 [=====] - 2s 27us/step - loss: 0.3681 - accuracy: 0.8991 - val_loss: 0.3471 - val_accuracy: 0.9064

Epoch 17/20

60000/60000 [=====] - 2s 31us/step - loss: 0.3640 - accuracy: 0.9005 - val_loss: 0.3438 - val_accuracy: 0.9070

Epoch 18/20

60000/60000 [=====] - 2s 28us/step - loss: 0.3603 - acc

uracy: 0.9012 - val_loss: 0.3404 - val_accuracy: 0.9073
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.3569 - acc
uracy: 0.9019 - val_loss: 0.3372 - val_accuracy: 0.9077
Epoch 20/20
60000/60000 [=====] - 2s 40us/step - loss: 0.3537 - acc
uracy: 0.9026 - val_loss: 0.3346 - val_accuracy: 0.9094

```

In [15]: score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.3346195185959339
Test accuracy: 0.9093999862670898



MLP + Sigmoid activation + SGDOptimizer

```
In [16]: # Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
In [17]: model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=[
'accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```


Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 9s 151us/step - loss: 2.2688 - accuracy: 0.2144 - val_loss: 2.2232 - val_accuracy: 0.4541

Epoch 2/20

60000/60000 [=====] - 10s 174us/step - loss: 2.1797 - accuracy: 0.4497 - val_loss: 2.1238 - val_accuracy: 0.4339

Epoch 3/20

60000/60000 [=====] - 9s 150us/step - loss: 2.0630 - accuracy: 0.5544 - val_loss: 1.9805 - val_accuracy: 0.6145

Epoch 4/20

60000/60000 [=====] - 9s 149us/step - loss: 1.8953 - accuracy: 0.6119 - val_loss: 1.7814 - val_accuracy: 0.6170

Epoch 5/20

60000/60000 [=====] - 8s 140us/step - loss: 1.6786 - accuracy: 0.6469 - val_loss: 1.5481 - val_accuracy: 0.6888

Epoch 6/20

60000/60000 [=====] - 7s 118us/step - loss: 1.4512 - accuracy: 0.6865 - val_loss: 1.3303 - val_accuracy: 0.7214

Epoch 7/20

60000/60000 [=====] - 8s 137us/step - loss: 1.2536 - accuracy: 0.7221 - val_loss: 1.1537 - val_accuracy: 0.7373

Epoch 8/20

60000/60000 [=====] - 10s 170us/step - loss: 1.0980 - accuracy: 0.7494 - val_loss: 1.0176 - val_accuracy: 0.7702

Epoch 9/20

60000/60000 [=====] - 8s 140us/step - loss: 0.9772 - accuracy: 0.7732 - val_loss: 0.9108 - val_accuracy: 0.7908

Epoch 10/20

60000/60000 [=====] - 12s 198us/step - loss: 0.8821 - accuracy: 0.7913 - val_loss: 0.8266 - val_accuracy: 0.8103

Epoch 11/20

60000/60000 [=====] - 13s 214us/step - loss: 0.8054 - accuracy: 0.8066 - val_loss: 0.7563 - val_accuracy: 0.8218

Epoch 12/20

60000/60000 [=====] - 11s 185us/step - loss: 0.7424 - accuracy: 0.8199 - val_loss: 0.6992 - val_accuracy: 0.8293

Epoch 13/20

60000/60000 [=====] - 15s 248us/step - loss: 0.6901 - accuracy: 0.8300 - val_loss: 0.6514 - val_accuracy: 0.8430

Epoch 14/20

60000/60000 [=====] - 9s 146us/step - loss: 0.6463 - accuracy: 0.8381 - val_loss: 0.6114 - val_accuracy: 0.8497

Epoch 15/20

60000/60000 [=====] - 8s 131us/step - loss: 0.6095 - accuracy: 0.8462 - val_loss: 0.5774 - val_accuracy: 0.8550

Epoch 16/20

60000/60000 [=====] - 9s 150us/step - loss: 0.5785 - accuracy: 0.8523 - val_loss: 0.5484 - val_accuracy: 0.8613

Epoch 17/20

60000/60000 [=====] - 8s 132us/step - loss: 0.5517 - accuracy: 0.8578 - val_loss: 0.5238 - val_accuracy: 0.8669

Epoch 18/20

60000/60000 [=====] - 9s 142us/step - loss: 0.5289 - accuracy: 0.8636 - val_loss: 0.5025 - val_accuracy: 0.8689

Epoch 19/20

60000/60000 [=====] - 7s 111us/step - loss: 0.5091 - accuracy: 0.8673 - val_loss: 0.4841 - val_accuracy: 0.8736

Epoch 20/20

60000/60000 [=====] - 8s 136us/step - loss: 0.4917 - accuracy: 0.8706 - val_loss: 0.4683 - val_accuracy: 0.8748

```

In [18]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.468309353351593
Test accuracy: 0.8748000264167786



```

In [20]: w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



```
In [0]: model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=
['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_ep
och, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290

Total params: 468,874
 Trainable params: 468,874
 Non-trainable params: 0

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
 60000/60000 [=====] - 3s 55us/step - loss: 0.5308 - acc: 0.8636 - val_loss: 0.2550 - val_acc: 0.9259
 Epoch 2/20
 53120/60000 [=====>....] - ETA: 0s - loss: 0.2244 - acc: 0.9340
 60000/60000 [=====] - 3s 51us/step - loss: 0.2205 - acc: 0.9351 - val_loss: 0.1946 - val_acc: 0.9417
 Epoch 3/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.1650 - acc: 0.9512 - val_loss: 0.1421 - val_acc: 0.9570
 Epoch 4/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.1279 - acc: 0.9614 - val_loss: 0.1238 - val_acc: 0.9645
 Epoch 5/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.1010 - acc: 0.9704 - val_loss: 0.1029 - val_acc: 0.9693
 Epoch 6/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.0801 - acc: 0.9763 - val_loss: 0.0877 - val_acc: 0.9725
 Epoch 7/20
 4480/60000 [=>.....] - ETA: 2s - loss: 0.0632 - acc: 0.9795
 60000/60000 [=====] - 3s 51us/step - loss: 0.0645 - acc: 0.9809 - val_loss: 0.0831 - val_acc: 0.9751
 Epoch 8/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.0519 - acc: 0.9842 - val_loss: 0.0724 - val_acc: 0.9780
 Epoch 9/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.0433 - acc: 0.9872 - val_loss: 0.0714 - val_acc: 0.9786
 Epoch 10/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.0347 - acc: 0.9898 - val_loss: 0.0695 - val_acc: 0.9776
 Epoch 11/20
 60000/60000 [=====] - 3s 51us/step - loss: 0.0268 - acc: 0.9930 - val_loss: 0.0659 - val_acc: 0.9796
 Epoch 12/20
 60000/60000 [=====] - 3s 52us/step - loss: 0.0219 - acc: 0.9944 - val_loss: 0.0642 - val_acc: 0.9809
 Epoch 13/20
 60000/60000 [=====] - 3s 50us/step - loss: 0.0180 - acc: 0.9953 - val_loss: 0.0677 - val_acc: 0.9794
 Epoch 14/20
 60000/60000 [=====] - 3s 50us/step - loss: 0.0133 - acc: 0.9970 - val_loss: 0.0647 - val_acc: 0.9803
 Epoch 15/20

60000/60000 [=====] - 3s 50us/step - loss: 0.0114 - acc: 0.9975 - val_loss: 0.0628 - val_acc: 0.9812
Epoch 16/20
57344/60000 [=====>..] - ETA: 0s - loss: 0.0085 - acc: 0.998260000/60000 [=====] - 3s 50us/step - loss: 0.0085 - acc: 0.9982 - val_loss: 0.0666 - val_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0070 - acc: 0.9986 - val_loss: 0.0643 - val_acc: 0.9822
Epoch 18/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0061 - acc: 0.9986 - val_loss: 0.0656 - val_acc: 0.9818
Epoch 19/20
60000/60000 [=====] - 3s 51us/step - loss: 0.0055 - acc: 0.9988 - val_loss: 0.0811 - val_acc: 0.9774
Epoch 20/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0038 - acc: 0.9992 - val_loss: 0.0723 - val_acc: 0.9818


```

In [0]: score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.06385514608082886

Test accuracy: 0.9824



```

In [0]: w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```

```

/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:588: FutureWarnin
g: remove_na is deprecated and is a private function. Do not use.
    kde_data = remove_na(group_data)
/usr/local/lib/python3.6/dist-packages/seaborn/categorical.py:816: FutureWarnin
g: remove_na is deprecated and is a private function. Do not use.
    violin_data = remove_na(group_data)

```

MLP + ReLU +SGD

```
In [0]: # Multilayer perceptron

# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i)}$ .
# h1 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0,\sigma) = N(0,0.062)$ 
# h2 =>  $\sigma=\sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0,\sigma) = N(0,0.125)$ 
# out =>  $\sigma=\sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290

Total params: 468,874
 Trainable params: 468,874
 Non-trainable params: 0

```
In [0]: model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 67us/step - loss: 0.7579 - acc: 0.7812 - val_loss: 0.3951 - val_acc: 0.8921

Epoch 2/20

60000/60000 [=====] - 4s 64us/step - loss: 0.3535 - acc: 0.8998 - val_loss: 0.3040 - val_acc: 0.9153

Epoch 3/20

60000/60000 [=====] - 4s 64us/step - loss: 0.2900 - acc: 0.9172 - val_loss: 0.2648 - val_acc: 0.9253

Epoch 4/20

60000/60000 [=====] - 4s 60us/step - loss: 0.2558 - acc: 0.9269 - val_loss: 0.2393 - val_acc: 0.9316

Epoch 5/20

60000/60000 [=====] - 4s 58us/step - loss: 0.2324 - acc: 0.9340 - val_loss: 0.2210 - val_acc: 0.9371

Epoch 6/20

60000/60000 [=====] - 4s 64us/step - loss: 0.2144 - acc: 0.9391 - val_loss: 0.2072 - val_acc: 0.9400

Epoch 7/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1995 - acc: 0.9443 - val_loss: 0.1957 - val_acc: 0.9444

Epoch 8/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1872 - acc: 0.9476 - val_loss: 0.1848 - val_acc: 0.9456

Epoch 9/20

60000/60000 [=====] - 3s 57us/step - loss: 0.1763 - acc: 0.9507 - val_loss: 0.1771 - val_acc: 0.9488

Epoch 10/20

60000/60000 [=====] - 4s 60us/step - loss: 0.1668 - acc: 0.9539 - val_loss: 0.1682 - val_acc: 0.9506

Epoch 11/20

60000/60000 [=====] - 4s 71us/step - loss: 0.1585 - acc: 0.9560 - val_loss: 0.1623 - val_acc: 0.9518

Epoch 12/20

60000/60000 [=====] - 7s 113us/step - loss: 0.1511 - acc: 0.9577 - val_loss: 0.1560 - val_acc: 0.9543

Epoch 13/20

60000/60000 [=====] - 7s 115us/step - loss: 0.1443 - acc: 0.9596 - val_loss: 0.1517 - val_acc: 0.9557

Epoch 14/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1379 - acc: 0.9615 - val_loss: 0.1474 - val_acc: 0.9572

Epoch 15/20

60000/60000 [=====] - 4s 66us/step - loss: 0.1323 - acc: 0.9628 - val_loss: 0.1429 - val_acc: 0.9580

Epoch 16/20

60000/60000 [=====] - 4s 69us/step - loss: 0.1270 - acc: 0.9645 - val_loss: 0.1371 - val_acc: 0.9598

Epoch 17/20

60000/60000 [=====] - 7s 110us/step - loss: 0.1221 - acc: 0.9661 - val_loss: 0.1351 - val_acc: 0.9602

Epoch 18/20

60000/60000 [=====] - 4s 62us/step - loss: 0.1177 - acc: 0.9671 - val_loss: 0.1309 - val_acc: 0.9618

Epoch 19/20

60000/60000 [=====] - 4s 60us/step - loss: 0.1136 - acc: 0.9685 - val_loss: 0.1263 - val_acc: 0.9631

Epoch 20/20

60000/60000 [=====] - 5s 79us/step - loss: 0.1094 - acc: 0.9694 - val_loss: 0.1241 - val_acc: 0.9631

```

In [0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```


Test score: 0.12405014228336513

Test accuracy: 0.9631



```

In [0]: w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



```
In [0]: model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290

=====
 Total params: 468,874
 Trainable params: 468,874
 Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 121us/step - loss: 0.2341 - acc: 0.9295 - val_loss: 0.1165 - val_acc: 0.9652

Epoch 2/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0878 - acc: 0.9729 - val_loss: 0.0883 - val_acc: 0.9720

Epoch 3/20

60000/60000 [=====] - 5s 75us/step - loss: 0.0544 - acc: 0.9825 - val_loss: 0.0860 - val_acc: 0.9729

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0354 - acc: 0.9885 - val_loss: 0.0699 - val_acc: 0.9797

Epoch 5/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0266 - acc: 0.9914 - val_loss: 0.0720 - val_acc: 0.9788

Epoch 6/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0200 - acc: 0.9941 - val_loss: 0.0696 - val_acc: 0.9803

Epoch 7/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0155 - acc: 0.9951 - val_loss: 0.0640 - val_acc: 0.9829

Epoch 8/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0140 - acc: 0.9952 - val_loss: 0.0848 - val_acc: 0.9792

Epoch 9/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0143 - acc: 0.9952 - val_loss: 0.0837 - val_acc: 0.9796

Epoch 10/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0128 - acc: 0.9958 - val_loss: 0.0946 - val_acc: 0.9782

Epoch 11/20

60000/60000 [=====] - 7s 125us/step - loss: 0.0081 - acc: 0.9974 - val_loss: 0.0682 - val_acc: 0.9826

Epoch 12/20

60000/60000 [=====] - 8s 129us/step - loss: 0.0121 - acc: 0.9959 - val_loss: 0.0793 - val_acc: 0.9816

Epoch 13/20

60000/60000 [=====] - 8s 133us/step - loss: 0.0107 - acc: 0.9963 - val_loss: 0.0746 - val_acc: 0.9820

Epoch 14/20

60000/60000 [=====] - 8s 129us/step - loss: 0.0113 - acc: 0.9960 - val_loss: 0.0813 - val_acc: 0.9816

Epoch 15/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0058 - acc:

c: 0.9982 - val_loss: 0.0770 - val_acc: 0.9842
Epoch 16/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0040 - ac
c: 0.9987 - val_loss: 0.0930 - val_acc: 0.9808
Epoch 17/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0119 - ac
c: 0.9959 - val_loss: 0.0813 - val_acc: 0.9819
Epoch 18/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0105 - ac
c: 0.9966 - val_loss: 0.1000 - val_acc: 0.9803
Epoch 19/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0064 - ac
c: 0.9981 - val_loss: 0.0852 - val_acc: 0.9831
Epoch 20/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0056 - ac
c: 0.9982 - val_loss: 0.1029 - val_acc: 0.9805

```

In [0]: score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10294274219236926

Test accuracy: 0.9805



```

In [0]: w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



MLP + Batch-Norm on hidden Layers + AdamOptimizer

```
In [24]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h3 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.055, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_6"

Layer (type)	Output Shape	Param #
dense_12 (Dense)	(None, 364)	285740
batch_normalization_5 (Batch Normalization)	(None, 364)	1456
dense_13 (Dense)	(None, 52)	18980
batch_normalization_6 (Batch Normalization)	(None, 52)	208
dense_14 (Dense)	(None, 10)	530
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		

```
In [25]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[  
        'accuracy'])  
  
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs,  
                           verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 12s 208us/step - loss: 0.3459 - accuracy: 0.8999 - val_loss: 0.2532 - val_accuracy: 0.9263

Epoch 2/20

60000/60000 [=====] - 12s 202us/step - loss: 0.2152 - accuracy: 0.9370 - val_loss: 0.2014 - val_accuracy: 0.9396

Epoch 3/20

60000/60000 [=====] - 10s 172us/step - loss: 0.1752 - accuracy: 0.9481 - val_loss: 0.1796 - val_accuracy: 0.9464

Epoch 4/20

60000/60000 [=====] - 10s 172us/step - loss: 0.1470 - accuracy: 0.9568 - val_loss: 0.1584 - val_accuracy: 0.9514

Epoch 5/20

60000/60000 [=====] - 10s 168us/step - loss: 0.1272 - accuracy: 0.9621 - val_loss: 0.1422 - val_accuracy: 0.9566

Epoch 6/20

60000/60000 [=====] - 10s 167us/step - loss: 0.1096 - accuracy: 0.9663 - val_loss: 0.1358 - val_accuracy: 0.9582

Epoch 7/20

60000/60000 [=====] - 10s 170us/step - loss: 0.0944 - accuracy: 0.9720 - val_loss: 0.1144 - val_accuracy: 0.9631

Epoch 8/20

60000/60000 [=====] - 10s 161us/step - loss: 0.0797 - accuracy: 0.9758 - val_loss: 0.1158 - val_accuracy: 0.9650

Epoch 9/20

60000/60000 [=====] - 9s 148us/step - loss: 0.0684 - accuracy: 0.9788 - val_loss: 0.1073 - val_accuracy: 0.9661

Epoch 10/20

60000/60000 [=====] - 10s 168us/step - loss: 0.0604 - accuracy: 0.9812 - val_loss: 0.1053 - val_accuracy: 0.9687

Epoch 11/20

60000/60000 [=====] - 9s 158us/step - loss: 0.0505 - accuracy: 0.9846 - val_loss: 0.1008 - val_accuracy: 0.9689

Epoch 12/20

60000/60000 [=====] - 10s 173us/step - loss: 0.0433 - accuracy: 0.9862 - val_loss: 0.1016 - val_accuracy: 0.9690

Epoch 13/20

60000/60000 [=====] - 12s 205us/step - loss: 0.0402 - accuracy: 0.9871 - val_loss: 0.1003 - val_accuracy: 0.9718

Epoch 14/20

60000/60000 [=====] - 10s 172us/step - loss: 0.0342 - accuracy: 0.9889 - val_loss: 0.1045 - val_accuracy: 0.9705

Epoch 15/20

60000/60000 [=====] - 11s 179us/step - loss: 0.0304 - accuracy: 0.9905 - val_loss: 0.0907 - val_accuracy: 0.9739

Epoch 16/20

60000/60000 [=====] - 10s 172us/step - loss: 0.0274 - accuracy: 0.9917 - val_loss: 0.0991 - val_accuracy: 0.9734

Epoch 17/20

60000/60000 [=====] - 10s 169us/step - loss: 0.0251 - accuracy: 0.9917 - val_loss: 0.1036 - val_accuracy: 0.9712

Epoch 18/20

60000/60000 [=====] - 10s 172us/step - loss: 0.0217 - accuracy: 0.9931 - val_loss: 0.1055 - val_accuracy: 0.9712

Epoch 19/20

60000/60000 [=====] - 12s 193us/step - loss: 0.0218 - accuracy: 0.9928 - val_loss: 0.0967 - val_accuracy: 0.9733

Epoch 20/20

60000/60000 [=====] - 11s 191us/step - loss: 0.0183 - a
ccuracy: 0.9942 - val_loss: 0.0927 - val_accuracy: 0.9770

```
In [0]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch,
# verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number
# of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10456635547156475

Test accuracy: 0.9732



```

In [0]: w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



5. MLP + Dropout + AdamOptimizer

```
In [23]: # https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras
```

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_10 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_11 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		


```
In [0]: model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 14s 227us/step - loss: 0.6612 - acc: 0.7951 - val_loss: 0.2860 - val_acc: 0.9166

Epoch 2/20

60000/60000 [=====] - 8s 136us/step - loss: 0.4250 - acc: 0.8710 - val_loss: 0.2545 - val_acc: 0.9252

Epoch 3/20

60000/60000 [=====] - 12s 198us/step - loss: 0.3841 - acc: 0.8846 - val_loss: 0.2391 - val_acc: 0.9298

Epoch 4/20

60000/60000 [=====] - 8s 138us/step - loss: 0.3551 - acc: 0.8927 - val_loss: 0.2279 - val_acc: 0.9325

Epoch 5/20

60000/60000 [=====] - 7s 123us/step - loss: 0.3355 - acc: 0.8986 - val_loss: 0.2127 - val_acc: 0.9356

Epoch 6/20

60000/60000 [=====] - 8s 136us/step - loss: 0.3234 - acc: 0.9031 - val_loss: 0.2029 - val_acc: 0.9387: 1s - loss:

Epoch 7/20

60000/60000 [=====] - 8s 131us/step - loss: 0.3068 - acc: 0.9077 - val_loss: 0.1927 - val_acc: 0.9421

Epoch 8/20

60000/60000 [=====] - 11s 185us/step - loss: 0.2933 - acc: 0.9113 - val_loss: 0.1836 - val_acc: 0.9453

Epoch 9/20

60000/60000 [=====] - 13s 222us/step - loss: 0.2850 - acc: 0.9131 - val_loss: 0.1797 - val_acc: 0.9451

Epoch 10/20

60000/60000 [=====] - 14s 236us/step - loss: 0.2715 - acc: 0.9187 - val_loss: 0.1738 - val_acc: 0.9465

Epoch 11/20

60000/60000 [=====] - 8s 141us/step - loss: 0.2611 - acc: 0.9214 - val_loss: 0.1671 - val_acc: 0.9506

Epoch 12/20

60000/60000 [=====] - 8s 134us/step - loss: 0.2464 - acc: 0.9252 - val_loss: 0.1554 - val_acc: 0.9525

Epoch 13/20

60000/60000 [=====] - 8s 137us/step - loss: 0.2382 - acc: 0.9278 - val_loss: 0.1479 - val_acc: 0.9554

Epoch 14/20

60000/60000 [=====] - 8s 136us/step - loss: 0.2275 - acc: 0.9313 - val_loss: 0.1375 - val_acc: 0.9580

Epoch 15/20

60000/60000 [=====] - 8s 137us/step - loss: 0.2183 - acc: 0.9337 - val_loss: 0.1326 - val_acc: 0.9599

Epoch 16/20

60000/60000 [=====] - 8s 138us/step - loss: 0.2068 - acc: 0.9384 - val_loss: 0.1297 - val_acc: 0.9613 loss: 0.2066 - acc

Epoch 17/20

60000/60000 [=====] - 8s 139us/step - loss: 0.2011 - acc: 0.9395 - val_loss: 0.1181 - val_acc: 0.9646

Epoch 18/20

60000/60000 [=====] - 8s 137us/step - loss: 0.1886 - acc: 0.9435 - val_loss: 0.1145 - val_acc: 0.9658

Epoch 19/20

60000/60000 [=====] - 8s 138us/step - loss: 0.1821 - acc: 0.9451 - val_loss: 0.1104 - val_acc: 0.9662

Epoch 20/20

```
60000/60000 [=====] - 8s 139us/step - loss: 0.1739 - ac  
c: 0.9473 - val_loss: 0.1093 - val_acc: 0.9679
```

```
In [0]: score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1093290721397847

Test accuracy: 0.9679



```

In [0]: w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



Hyper-parameter tuning of Keras models using Sklearn

```
In [0]: from keras.optimizers import Adam,RMSprop,SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

```
In [0]: # https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid','relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)
```

```
In [0]: print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

Best: 0.975633 using {'activ': 'relu'}
0.974650 (0.001138) with: {'activ': 'sigmoid'}
0.975633 (0.002812) with: {'activ': 'relu'}
```

Additional Tries from Assignments

Model 1: MLP + Batch-Norm on hidden Layers + AdamOptimizer change dense of hidden layers </2>

```
In [28]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 364)	285740
batch_normalization_7 (Batch Normalization)	(None, 364)	1456
dense_16 (Dense)	(None, 52)	18980
batch_normalization_8 (Batch Normalization)	(None, 52)	208
dense_17 (Dense)	(None, 10)	530
Total params: 306,914		
Trainable params: 306,082		
Non-trainable params: 832		


```
In [30]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[  
        'accuracy'])  
  
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs,  
                           verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 180us/step - loss: 0.3523 - accuracy: 0.8970 - val_loss: 0.2629 - val_accuracy: 0.9195

Epoch 2/20

60000/60000 [=====] - 8s 139us/step - loss: 0.2246 - accuracy: 0.9340 - val_loss: 0.2214 - val_accuracy: 0.9339 loss: 0.2247 - accuracy: 0.

Epoch 3/20

60000/60000 [=====] - 8s 129us/step - loss: 0.1836 - accuracy: 0.9452 - val_loss: 0.1891 - val_accuracy: 0.9442

Epoch 4/20

60000/60000 [=====] - 8s 135us/step - loss: 0.1511 - accuracy: 0.9553 - val_loss: 0.1742 - val_accuracy: 0.9493

Epoch 5/20

60000/60000 [=====] - 8s 129us/step - loss: 0.1325 - accuracy: 0.9612 - val_loss: 0.1527 - val_accuracy: 0.9564

Epoch 6/20

60000/60000 [=====] - 8s 132us/step - loss: 0.1122 - accuracy: 0.9664 - val_loss: 0.1412 - val_accuracy: 0.9576

Epoch 7/20

60000/60000 [=====] - 9s 151us/step - loss: 0.0967 - accuracy: 0.9706 - val_loss: 0.1268 - val_accuracy: 0.9637

Epoch 8/20

60000/60000 [=====] - 9s 158us/step - loss: 0.0832 - accuracy: 0.9747 - val_loss: 0.1251 - val_accuracy: 0.9635

Epoch 9/20

60000/60000 [=====] - 9s 145us/step - loss: 0.0711 - accuracy: 0.9790 - val_loss: 0.1152 - val_accuracy: 0.9664

Epoch 10/20

60000/60000 [=====] - 9s 155us/step - loss: 0.0609 - accuracy: 0.9814 - val_loss: 0.1119 - val_accuracy: 0.9683

Epoch 11/20

60000/60000 [=====] - 8s 140us/step - loss: 0.0543 - accuracy: 0.9829 - val_loss: 0.1120 - val_accuracy: 0.9682

Epoch 12/20

60000/60000 [=====] - 14s 237us/step - loss: 0.0455 - accuracy: 0.9858 - val_loss: 0.1063 - val_accuracy: 0.9678

Epoch 13/20

60000/60000 [=====] - 12s 198us/step - loss: 0.0396 - accuracy: 0.9873 - val_loss: 0.1053 - val_accuracy: 0.9715

Epoch 14/20

60000/60000 [=====] - 13s 213us/step - loss: 0.0341 - accuracy: 0.9895 - val_loss: 0.1029 - val_accuracy: 0.9707

Epoch 15/20

60000/60000 [=====] - 10s 173us/step - loss: 0.0318 - accuracy: 0.9899 - val_loss: 0.0992 - val_accuracy: 0.9726

Epoch 16/20

60000/60000 [=====] - 10s 170us/step - loss: 0.0268 - accuracy: 0.9915 - val_loss: 0.1039 - val_accuracy: 0.9732

Epoch 17/20

60000/60000 [=====] - 10s 169us/step - loss: 0.0250 - accuracy: 0.9924 - val_loss: 0.1096 - val_accuracy: 0.9718

Epoch 18/20

60000/60000 [=====] - 10s 166us/step - loss: 0.0221 - accuracy: 0.9931 - val_loss: 0.1113 - val_accuracy: 0.9713

Epoch 19/20

60000/60000 [=====] - 10s 172us/step - loss: 0.0205 - accuracy: 0.9933 - val_loss: 0.1071 - val_accuracy: 0.9725

Epoch 20/20
60000/60000 [=====] - 11s 180us/step - loss: 0.0171 - a
ccuracy: 0.9948 - val_loss: 0.1001 - val_accuracy: 0.9752

Before (hidden 1 = 512, hidden 2 = 128)

```
In [0]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10456635547156475

Test accuracy: 0.9732



After (hidden 1 = 364, hidden 2 = 52)

```
In [31]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1000667967731366
Test accuracy: 0.9751999974250793



```

In [32]: w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```



conclusion for model 1: decrease activation units from 512 to 264 did not impact result much, where accuray or loss did not decrease

In []:

Model 2: MLP + Batch-Norm on hidden Layers + AdamOptimizer Add + 3 hidden layer </2>

```
In [33]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(364, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_18 (Dense)	(None, 364)	285740
batch_normalization_9 (Batch Normalization)	(None, 364)	1456
dense_19 (Dense)	(None, 128)	46720
batch_normalization_10 (Batch Normalization)	(None, 128)	512
dense_20 (Dense)	(None, 52)	6708
batch_normalization_11 (Batch Normalization)	(None, 52)	208
dense_21 (Dense)	(None, 10)	530
Total params: 341,874		
Trainable params: 340,786		
Non-trainable params: 1,088		

```
In [34]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[
          'accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs,
                           verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 13s 213us/step - loss: 0.3049 - accuracy: 0.9117 - val_loss: 0.2081 - val_accuracy: 0.9423

Epoch 2/20

60000/60000 [=====] - 12s 192us/step - loss: 0.1666 - accuracy: 0.9516 - val_loss: 0.1531 - val_accuracy: 0.9548

Epoch 3/20

60000/60000 [=====] - 11s 189us/step - loss: 0.1254 - accuracy: 0.9628 - val_loss: 0.1253 - val_accuracy: 0.9629

Epoch 4/20

60000/60000 [=====] - 10s 164us/step - loss: 0.0976 - accuracy: 0.9701 - val_loss: 0.1188 - val_accuracy: 0.9637

Epoch 5/20

60000/60000 [=====] - 10s 168us/step - loss: 0.0788 - accuracy: 0.9758 - val_loss: 0.1151 - val_accuracy: 0.9658

Epoch 6/20

60000/60000 [=====] - 11s 179us/step - loss: 0.0673 - accuracy: 0.9791 - val_loss: 0.1012 - val_accuracy: 0.9703

Epoch 7/20

60000/60000 [=====] - 11s 182us/step - loss: 0.0546 - accuracy: 0.9834 - val_loss: 0.0997 - val_accuracy: 0.9684

Epoch 8/20

60000/60000 [=====] - 11s 180us/step - loss: 0.0470 - accuracy: 0.9852 - val_loss: 0.0979 - val_accuracy: 0.9724

Epoch 9/20

60000/60000 [=====] - 11s 180us/step - loss: 0.0385 - accuracy: 0.9874 - val_loss: 0.0903 - val_accuracy: 0.9730

Epoch 10/20

60000/60000 [=====] - 11s 179us/step - loss: 0.0328 - accuracy: 0.9894 - val_loss: 0.0875 - val_accuracy: 0.9744

Epoch 11/20

60000/60000 [=====] - 10s 173us/step - loss: 0.0275 - accuracy: 0.9908 - val_loss: 0.0889 - val_accuracy: 0.9757

Epoch 12/20

60000/60000 [=====] - 11s 185us/step - loss: 0.0277 - accuracy: 0.9911 - val_loss: 0.0892 - val_accuracy: 0.9745

Epoch 13/20

60000/60000 [=====] - 11s 187us/step - loss: 0.0234 - accuracy: 0.9926 - val_loss: 0.0909 - val_accuracy: 0.9740

Epoch 14/20

60000/60000 [=====] - 10s 163us/step - loss: 0.0195 - accuracy: 0.9936 - val_loss: 0.0939 - val_accuracy: 0.9750

Epoch 15/20

60000/60000 [=====] - 11s 182us/step - loss: 0.0186 - accuracy: 0.9940 - val_loss: 0.0921 - val_accuracy: 0.9760

Epoch 16/20

60000/60000 [=====] - 11s 186us/step - loss: 0.0175 - accuracy: 0.9943 - val_loss: 0.0944 - val_accuracy: 0.9759

Epoch 17/20

60000/60000 [=====] - 10s 168us/step - loss: 0.0142 - accuracy: 0.9953 - val_loss: 0.0895 - val_accuracy: 0.9763

Epoch 18/20

60000/60000 [=====] - 11s 187us/step - loss: 0.0157 - accuracy: 0.9950 - val_loss: 0.0888 - val_accuracy: 0.9772

Epoch 19/20

60000/60000 [=====] - 12s 200us/step - loss: 0.0160 - accuracy: 0.9947 - val_loss: 0.1010 - val_accuracy: 0.9740

Epoch 20/20

60000/60000 [=====] - 12s 202us/step - loss: 0.0147 - a
ccuracy: 0.9952 - val_loss: 0.0932 - val_accuracy: 0.9765

```
In [35]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0932293744858849
Test accuracy: 0.9764999747276306



conclusion for model 2: Adding another hidden layer slightly increased accuracy, and reduced log loss

In []:

Model 3: MLP + Batch-Norm on hidden Layers + AdamOptimizer Add + 5 hidden layer

```

In [38]: # Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma=\sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma=\sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(364, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(32, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```


Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 512)	401920
batch_normalization_12 (Batch Normalization)	(None, 512)	2048
dense_23 (Dense)	(None, 364)	186732
batch_normalization_13 (Batch Normalization)	(None, 364)	1456
dense_24 (Dense)	(None, 128)	46720
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dense_25 (Dense)	(None, 52)	6708
batch_normalization_15 (Batch Normalization)	(None, 52)	208
dense_26 (Dense)	(None, 32)	1696
batch_normalization_16 (Batch Normalization)	(None, 32)	128
dense_27 (Dense)	(None, 10)	330
Total params: 648,458		
Trainable params: 646,282		
Non-trainable params: 2,176		

```
In [39]: model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=[
        'accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epochs,
                           verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 20s 328us/step - loss: 0.3063 - accuracy: 0.9132 - val_loss: 0.2306 - val_accuracy: 0.9405

Epoch 2/20

60000/60000 [=====] - 19s 318us/step - loss: 0.1589 - accuracy: 0.9525 - val_loss: 0.1302 - val_accuracy: 0.9620

Epoch 3/20

60000/60000 [=====] - 20s 331us/step - loss: 0.1162 - accuracy: 0.9648 - val_loss: 0.1165 - val_accuracy: 0.9654

Epoch 4/20

60000/60000 [=====] - 19s 320us/step - loss: 0.0910 - accuracy: 0.9722 - val_loss: 0.1206 - val_accuracy: 0.9646

Epoch 5/20

60000/60000 [=====] - 20s 329us/step - loss: 0.0758 - accuracy: 0.9766 - val_loss: 0.0935 - val_accuracy: 0.9736

Epoch 6/20

60000/60000 [=====] - 18s 300us/step - loss: 0.0611 - accuracy: 0.9814 - val_loss: 0.1113 - val_accuracy: 0.9665

Epoch 7/20

60000/60000 [=====] - 21s 356us/step - loss: 0.0514 - accuracy: 0.9838 - val_loss: 0.0948 - val_accuracy: 0.9718

Epoch 8/20

60000/60000 [=====] - 28s 466us/step - loss: 0.0464 - accuracy: 0.9851 - val_loss: 0.0873 - val_accuracy: 0.9739

Epoch 9/20

60000/60000 [=====] - 28s 464us/step - loss: 0.0403 - accuracy: 0.9871 - val_loss: 0.1081 - val_accuracy: 0.9694

Epoch 10/20

60000/60000 [=====] - 28s 466us/step - loss: 0.0343 - accuracy: 0.9883 - val_loss: 0.0973 - val_accuracy: 0.9733

Epoch 11/20

60000/60000 [=====] - 26s 429us/step - loss: 0.0353 - accuracy: 0.9889 - val_loss: 0.0992 - val_accuracy: 0.9724

Epoch 12/20

60000/60000 [=====] - 24s 400us/step - loss: 0.0259 - accuracy: 0.9917 - val_loss: 0.0981 - val_accuracy: 0.9743

Epoch 13/20

60000/60000 [=====] - 27s 450us/step - loss: 0.0257 - accuracy: 0.9920 - val_loss: 0.1093 - val_accuracy: 0.9701

Epoch 14/20

60000/60000 [=====] - 26s 427us/step - loss: 0.0253 - accuracy: 0.9915 - val_loss: 0.0944 - val_accuracy: 0.9752

Epoch 15/20

60000/60000 [=====] - 24s 395us/step - loss: 0.0209 - accuracy: 0.9931 - val_loss: 0.0922 - val_accuracy: 0.9766

Epoch 16/20

60000/60000 [=====] - 20s 328us/step - loss: 0.0225 - accuracy: 0.9928 - val_loss: 0.1058 - val_accuracy: 0.9731

Epoch 17/20

60000/60000 [=====] - 20s 335us/step - loss: 0.0223 - accuracy: 0.9927 - val_loss: 0.0872 - val_accuracy: 0.9768

Epoch 18/20

60000/60000 [=====] - 20s 339us/step - loss: 0.0161 - accuracy: 0.9950 - val_loss: 0.0964 - val_accuracy: 0.9773

Epoch 19/20

60000/60000 [=====] - 18s 305us/step - loss: 0.0200 - accuracy: 0.9934 - val_loss: 0.1071 - val_accuracy: 0.9758

Epoch 20/20

60000/60000 [=====] - 20s 332us/step - loss: 0.0167 - a
ccuracy: 0.9942 - val_loss: 0.0932 - val_accuracy: 0.9780

```
In [40]: score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09317940528480685
Test accuracy: 0.9779999852180481



conclusion for model 3: Adding two more hidden layer slightly increased accuracy, and reduced log loss

In []: